

Sasank Jampana

📞 +91 9347004843 ✉ sasankjampana8@gmail.com 🔗 [linkedin.com/sasankjampana](https://www.linkedin.com/sasankjampana) 🐙 github.com/sasankjampana8

Summary

- Full-stack Data Scientist with 2 years of experience in developing and deploying end-to-end machine learning solutions. Expertise in core data science, cloud computing, and scalable model deployment. Passionate about leveraging Generative AI, LLMs, and MLOps to build intelligent, production-ready applications.

Education

Vellore Institute of Technology, Andhra Pradesh

Bachelor of Technology in Computer Science (CGPA: 8.93 / 10.00)

July 2019 - May 2023

Amaravati, AP

Experience

Intellect Design Arena

Data Scientist

June 2023 – Present

Hyderabad, Telangana

- Engineered an ML management framework to streamline model development, training, deployment, and monitoring for data scientists, designed for seamless on-premises deployment and scalable lifecycle management.
- Orchestrated a ReAct-based LLM chatbot integrating a RAG-powered knowledge retrieval system (leveraging Claude Sonnet 3.5), a Text-to-SQL agent for PostgreSQL query generation (powered by OpenAI GPT-4o), a FastAPI backend for secure PII-masked data retrieval, an S3-backed visualization tool, and a robust dataset lifecycle and knowledge base management system.
- Developed a regex-based document extraction system for CAS parsing, supporting KFinTech, NSDL, and CDSL documents with 99% accuracy and extraction time under 3 seconds.
- Designed and deployed a high-performance stacked ensemble multi-label classification model for 10 categories, achieving a 96.2% macro F1-score. Enhanced model interpretability using XAI techniques and implemented drift detection mechanisms. Deployed on AWS and on-premises using Docker and Flask for real-time and batch inference.
- Constructed a binary classification model with a 99% weighted F1-score, integrating drift detection and explainability techniques. Built end-to-end training and inference pipelines leveraging Docker and SageMaker for scalable deployment.
- Architected an Intelligent Document Processing (IDP) system utilizing Transformers, Word2Vec, Pytesseract, BERT, and LayoutLM V1 for automated document extraction, categorization, and clustering, optimizing document workflows with minimal human intervention.

Forbes Advisor

Data Engineer Intern

Dec 2022 – May 2023

Chennai, Tamil Nadu

- Developed a proof of concept (PoC) for real-time data pipelines on Google Cloud Platform (GCP), integrating diverse data sources and validating data before ingestion into Google BigQuery, enhancing data processing efficiency and analytical performance.

Technical Skills

Programming & Scripting: Python, MySQL

Machine Learning & AI: TensorFlow, PyTorch, Scikit-Learn, Hugging Face

Generative AI: LangChain, LangGraph, Prompt Engineering, Fine-tuning, Vector DBs

Cloud & MLOps: AWS (SageMaker, Lambda, S3), GCP (BigQuery), Docker, Kubernetes, MLflow, FastAPI and Flask

Databases: PostgreSQL, MySQL, MongoDB, Pyspark

Certifications & Achievements

Diploma in Data Science, IIT Madras

Spotlight Award: Recognized for exceptional contributions to the Customer Engagement Indicator project, acknowledged as a valuable resource for driving the project's success and delivering impactful results.